

Assignment 4 Report

In this assignment, I have tried the whole dataset in using `sklearn.dataset.fetch_20newsgroups`

File structure

```
.
├── checkpoints: ignored by git
├── models
├── preprocess.py: fetch dataset from sklearn and pass it to dataSet of fastNLP
├── test_cnn.py: reload saved model and test accuracy
├── train_cnn.py
└── train_rnn.py
```

Usage

Train RNN model:

```
python3 train_rnn.py
```

Train CNN model:

```
python3 run_cnn.py
```

RNN Text Classification

In general, the classification is mainly done by the last layer's hidden state of LSTM after feeding the sentence (article).

The model is

- Embedding: `(seq_len, batch_size) -> (seq_len, batch_size, embedding_dim)`
- LSTM: `(seq_len, batch_size, embedding_dim) -> hidden:(num_layer*num_directions, batch_size, hidden_size)`
 - `num_directions` is 2, a bidirectional LSTM
- Concatenation (of hidden state of last layer of LSTM): `(batch_size, hidden_size) -> (batch_size, hidden_size*2)`
- Dropout: input shape is the same as output shape
- Linear: `(batch_size, hidden_size*2) -> (batch_size, output_size=num_classes)`

CNN Text Classification

My CNN model is just a simple reproduction of the paper(refer to [implementation from fastNLP](#)). CNN is much faster than RNN since it can make full use of CPU.

The CNN model for text classification is as follows:

- Embedding: `(batch_size, seq_len) -> (batch_size, seq_len, embedding_dim)`
- Convolution: `(batch_size, height_in=seq_len, width_in=embedding_dim) -> (kernel_types, batch_size, out_channels, height_out, width_out=1)`
 - `kernel_types`: 3 kinds of kernels with height 3, 4, 5. I.e. 3-gram, 4-gram and 5-gram.
 - `out_channels`: 100 kernels of each kind.
- Relu: input shape is the same as output shape
- Max-pooling: `(kernel_types, batch_size, out_channels, height_out) -> (kernel_types, batch_size, out_channels)`
- Concatenation: `(kernel_types, batch_size, out_channels) -> (batch_size, kernel_types*out_channels)`
- Dropout: input shape is the same as output shape
- Linear: `(batch_size, kernel_types*out_channels) -> (batch_size, output_size=num_classes)`

After one night, **the accuracy on test set is 0.69**, and I think it may due to the large dataset, which make it unable to converge in several hours.

Suggestions for fastNLP

When getting familiar with fastNLP, I find `DataSet` really a handy one and happily get rid of handwriting trivial preprocessing data. Still, something can improve to make it more efficient and user-friendly.

- More on padding sequence:

It seems that padded data provided by `DataSet` just simply pad some 0 over the variant length sequences. Supports are recommended to make it more memory efficient by **using `pad_sequence` in Pytorch**. Then the inbuilt `rnn` in pytorch will handle our padded sequence quite well (This requires add some wrapper around `rnn`, which is shown in `./models/custom_rnn.py`). That's padding value in the padded sequence padded by `pad_sequence` won't be taken into consideration in RNN, which can greatly improve efficiency with the sequences' length varies differently.

- Save/load model and optimizer:

I'm afraid that the saving of optimizer hasn't been implemented, which is required for optimizers with inner adaptive parameters such as Adam. And it would be much better if `Trainer` can automatically (or ask the user to determine) **restart from the previously saved checkpoint** if exists (then I don't have to code for loading the model and optimizer).

